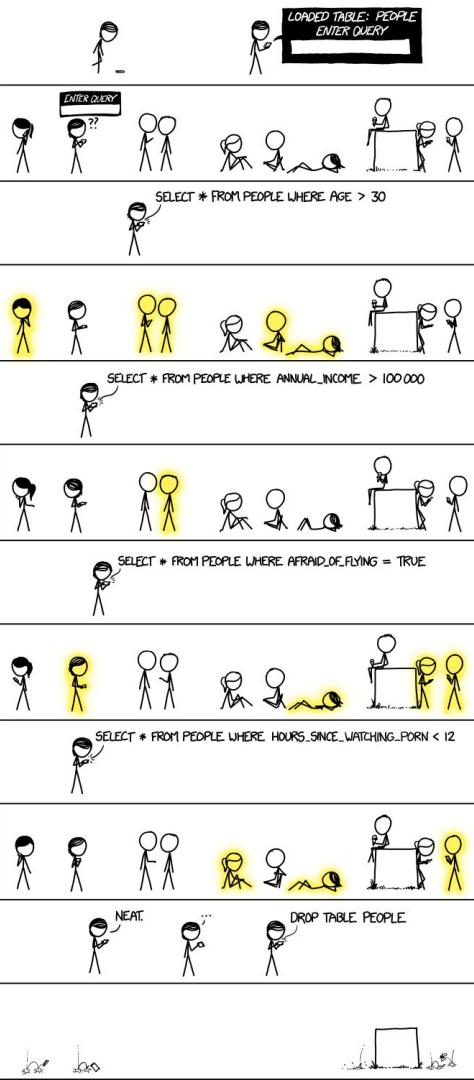


Session 7

Relational Databases (SQL II)



Remember back
9,900 minutes ago...

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68
S104	IAI	65
S106	PR2	43
S107	PR1	76
S107	PR2	60
S107	IAI	35

- **CREATE TABLE** - creates a table
- **SELECT** - return the specified column(s) from a table
- **UPDATE** - change row(s) in a table
- **WHERE** - filters row(s) by condition
- **INSERT** - add a row to a table
- **DELETE** - remove row(s) from a table

All operate on a single table.

All return a single table.

Think in tables, entities & attributes...



Remember back
9,900 minutes ago...



Remember back
9,900 minutes ago...

student

id	first	last
S103	John	Smith
S104	Mary	Jones
S105	Jane	Brown
S106	Mark	Jones
S107	John	Brown

Only ask about students

Before:

Operated on **one table**.

Only could ask questions about one set
of entities.



Remember back
9,900 minutes ago...

student

id	first	last
S103	John	Smith
S104	Mary	Jones
S105	Jane	Brown
S106	Mark	Jones

Only ask about students

course

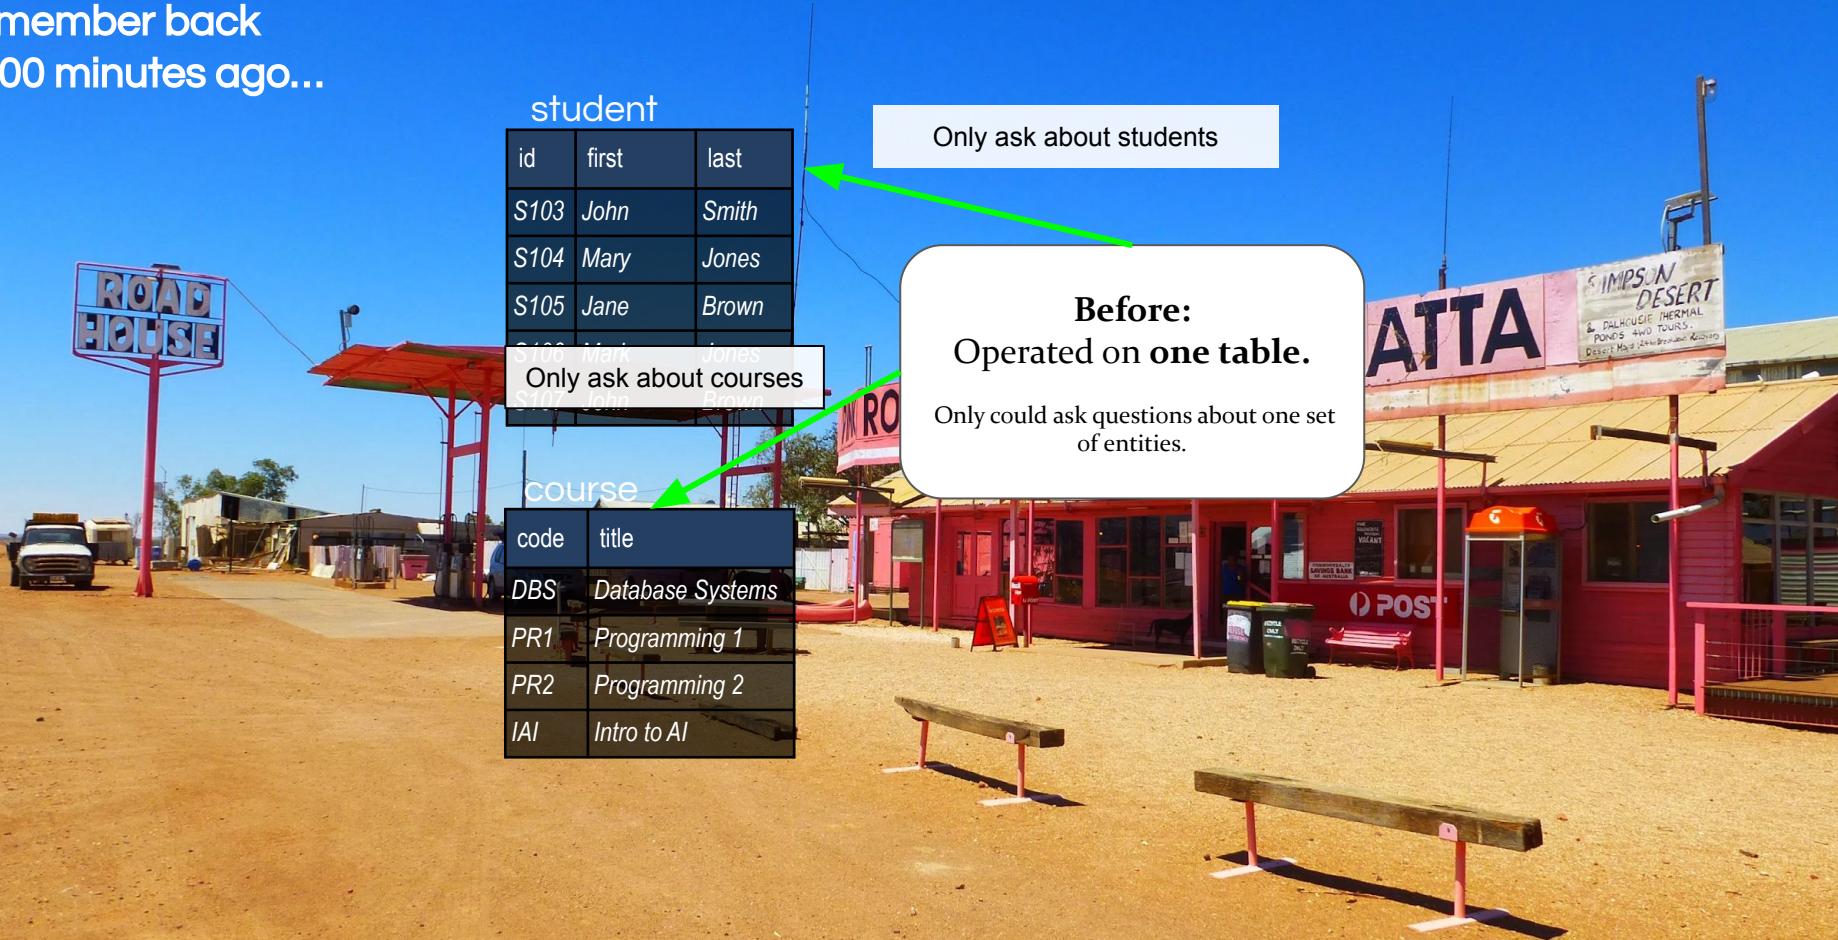
code	title
DBS	Database Systems
PR1	Programming 1
PR2	Programming 2
IAI	Intro to AI

Only ask about courses

Before:

Operated on one table.

Only could ask questions about one set of entities.



Remember back
9,900 minutes ago...

grade

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68
S104	IAI	65
S106	PR2	43
S107	PR1	76
S107	PR2	60
S107	IAI	35

Only ask about grades

student

id	first	last
S103	John	Smith
S104	Mary	Jones
S105	Jane	Brown

S106 Mark Jones
Only ask about courses

course

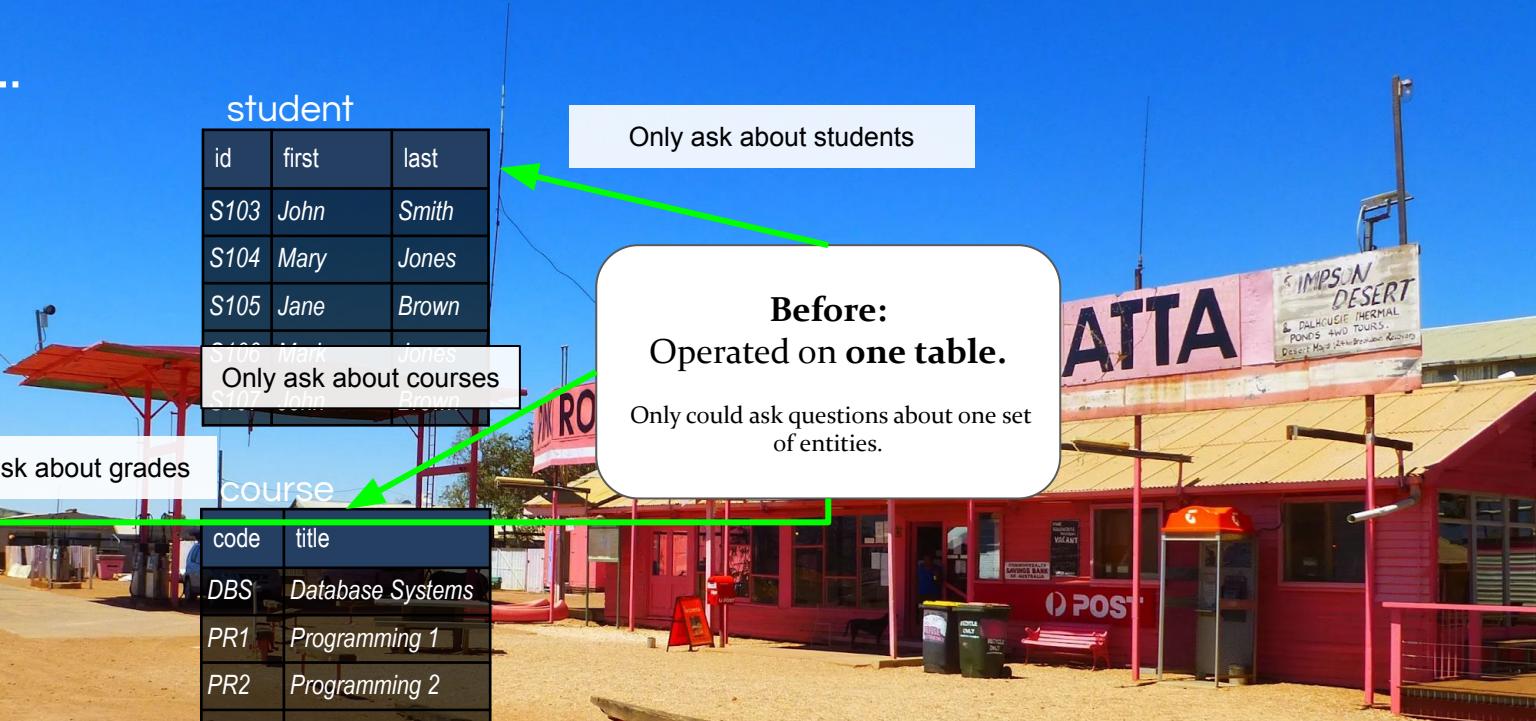
code	title
DBS	Database Systems
PR1	Programming 1
PR2	Programming 2
IAI	Intro to AI

Only ask about students

Before:

Operated on one table.

Only could ask questions about one set of entities.



grade

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68
S104	IAI	65
S106	PR2	43
S107	PR1	76
S107	PR2	60
S107	IAI	35

student

id	first	last
S103	John	Smith
S104	Mary	Jones
S105	Jane	Brown
S106	Mark	Jones
S107	John	Brown

course

code	title
DBS	Database Systems
PR1	Programming 1
PR2	Programming 2
IAI	Intro to AI

Now: Joining tables together.

Run operators (ask questions) on new set of composite entities.



SQL: JOINING TABLES

(the traditional way)

grade

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68
S104	IAI	65
S106	PR2	43
S107	PR1	76
S107	PR2	60
S107	IAI	35

student

id	first	last
S103	John	Smith
S104	Mary	Jones
S105	Jane	Brown
S106	Mark	Jones
S107	John	Brown

course

code	title
DBS	Database Systems
PR1	Programming 1
PR2	Programming 2
IAI	Intro to AI

Facts about two different entities types can be joined together to obtain facts about pairs of entity occurrences (one from each type).

We're going to learn the mechanism first, then interpretation.

Joining student occurrence: S103, John, Smith

With grade occurrence: S103, DBS, 72

Creates a new "Graded student" entity. (more on interpretation later)

Joining grade occurrence: S103, DBS, 72

With course occurrence: DBS, Database Systems

SQL: JOINING TABLES

(the traditional way)

grade

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68
S104	IAI	65
S106	PR2	43
S107	PR1	76
S107	PR2	60
S107	IAI	35

student

id	first	last
S103	John	Smith
S104	Mary	Jones
S105	Jane	Brown
S106	Mark	Jones
S107	John	Brown

course

code	title
DBS	Database Systems
PR1	Programming 1
PR2	Programming 2
IAI	Intro to AI

Facts about two different entities types can be joined together to obtain facts about pairs of entity occurrences (one from each type).

We're going to learn the mechanism first, then interpretation.

Joining student occurrence: S103, John, Smith

With grade occurrence: S103, DBS, 72

Creates a new "Graded student" entity. (more on interpretation later)

Joining grade occurrence: S103, DBS, 72

With course occurrence: DBS, Database Systems

One-to-many relationship:

student 1...1 → has → 0..* grades

→ Grade table is appended / extended with its unique match in student

→ Creates a new, more detailed, grade entity.

SQL: JOINING TABLES

(the traditional way)

student

id	first	last
S103	John	Smith
S104	Mary	Jones
S105	Jane	Brown
S106	Mark	Jones
S107	John	Brown

id	first	last	id	code	mark
S103	John	Smith	S103	DBS	72
			S103	IAI	58
			S104	PR1	68
			S104	IAI	65
			S106	PR2	43
			S107	PR1	76
			S107	PR2	60
			S107	IAI	35
S104	Mary	Jones	S103	DBS	72
			S103	IAI	58
			S104	PR1	68
			S104	IAI	65
			S106	PR2	43
			S107	PR1	76
			S107	PR2	60
			S107	IAI	35
S105	Jane	Brown	S103	DBS	72
			S103	IAI	58
			S104	PR1	68
			S104	IAI	65
			S106	PR2	43
			S107	PR1	76
			S107	PR2	60
			S107	IAI	35
			S104	IAI	65
			...		

Joining tables joins **every row** in the first table to **every row** in the second.

grade

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68
S104	IAI	65
S106	PR2	43
S107	PR1	76
S107	PR2	60
S107	IAI	35

```
SELECT *  
FROM student, grade
```

SQL: JOINING TABLES

(the traditional way)

student

id	first	last
S103	John	Smith
S104	Mary	Jones
S105	Jane	Brown
S106	Mark	Jones
S107	John	Brown

id	first	last	id	code	mark
S103	John	Smith	S103	DBS	72
			S103	IAI	58
			S104	PR1	68
			S104	IAI	65
			S106	PR2	43
			S107	PR1	76
			S107	PR2	60
			S107	IAI	35
S104	Mary	Jones	S103	DBS	72
			S103	IAI	58
			S104	PR1	68
			S104	IAI	65
			S106	PR2	43
			S107	PR1	76
			S107	PR2	60
			S107	IAI	35
S105	Jane	Brown	S103	DBS	72
			S103	IAI	58
			S104	PR1	68
			S104	IAI	65
			S106	PR2	43
			S107	PR1	76
			S107	PR2	60
			S107	IAI	35

Joining tables joins **every row** in the first table to **every row** in the second.

Database does not understand what resultant rows are **true and should be kept**.

```
SELECT *  
FROM student, grade
```

Database does not understand
what resultant rows are **true and**
should be kept.

YOU must tell it.

student

id	first	last
S103	John	Smith
S104	Mary	Jones
S105	Jane	Brown
S106	Mark	Jones
S107	John	Brown

grade

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68
S104	IAI	65
S106	PR2	43
S107	PR1	76
S107	PR2	60
S107	IAI	35

id	first	last	id	code	mark
S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58
S103	John	Smith	S104	PR1	68
S103	John	Smith	S104	IAI	65
S103	John	Smith	S106	PR2	43
S103	John	Smith	S107	PR1	76
S103	John	Smith	S107	PR2	60
S103	John	Smith	S107	IAI	35
S104	Mary	Jones	S103	DBS	72
S104	Mary	Jones	S103	IAI	58
S104	Mary	Jones	S104	PR1	68
S104	Mary	Jones	S104	IAI	65
S104	Mary	Jones	S106	PR2	43
S104	Mary	Jones	S107	PR1	76
S104	Mary	Jones	S107	PR2	60
S104	Mary	Jones	S107	IAI	35
S105	Jane	Brown	S103	DBS	72
S105	Jane	Brown	S103	IAI	58
S105	Jane	Brown	S104	PR1	68
S105	Jane	Brown	S104	IAI	65
S105	Jane	Brown	S106	PR2	43
S105	Jane	Brown	S107	PR1	76
S105	Jane	Brown	S107	PR2	60
S105	Jane	Brown	S107	IAI	35

```
SELECT *
FROM student, grade
WHERE student.id = grade.id
```

Database does not understand what resultant rows are **true and should be kept**.

YOU must tell it.

student

id	first	last
S103	John	Smith
S104	Mary	Jones
S105	Jane	Brown
S106	Mark	Jones
S107	John	Brown

grade

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68
S104	IAI	65
S106	PR2	43
S107	PR1	76
S107	PR2	60
S107	IAI	35

id	first	last	id	code	mark
S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58
S103	John	Smith	S104	PR1	68
S103	John	Smith	S104	IAI	65
S103	John	Smith	S106	PR2	43
S103	John	Smith	S107	PR1	76
S103	John	Smith	S107	PR2	60
S103	John	Smith	S107	IAI	35
S104	Mary	Jones	S103	DBS	72
S104	Mary	Jones	S103	IAI	58
S104	Mary	Jones	S104	PR1	68
S104	Mary	Jones	S104	IAI	65
S104	Mary	Jones	S106	PR2	43
S104	Mary	Jones	S107	PR1	76
S104	Mary	Jones	S107	PR2	60
S104	Mary	Jones	S107	IAI	35
S105	Jane	Brown	S103	DBS	72
S105	Jane	Brown	S103	IAI	58
S105	Jane	Brown	S104	PR1	68
S105	Jane	Brown	S104	IAI	65
S105	Jane	Brown	S106	PR2	43
S105	Jane	Brown	S107	PR1	76
S105	Jane	Brown	S107	PR2	60
S105	Jane	Brown	S107	IAI	35

```
SELECT *  
FROM student, grade  
WHERE student.id = grade.id
```

Database does not understand what resultant rows are **true and should be kept**.

YOU must tell it.

student

id	first	last
S103	John	Smith
S104	Mary	Jones
S105	Jane	Brown
S106	Mark	Jones
S107	John	Brown

grade

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68
S104	IAI	65
S106	PR2	43
S107	PR1	76
S107	PR2	60
S107	IAI	35

id	first	last	id	code	mark
S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58
S103	John	Smith	S104	PR1	68
S103	John	Smith	S104	IAI	65
S103	John	Smith	S106	PR2	43
S103	John	Smith	S107	PR1	76
S103	John	Smith	S107	PR2	60
S103	John	Smith	S107	IAI	35
S104	Mary	Jones	S103	DBS	72
S104	Mary	Jones	S103	IAI	58
S104	Mary	Jones	S104	PR1	68
S104	Mary	Jones	S104	IAI	65
S104	Mary	Jones	S106	PR2	43
S104	Mary	Jones	S107	PR1	76
S104	Mary	Jones	S107	PR2	60
S104	Mary	Jones	S107	IAI	35
S105	Jane	Brown	S103	DBS	72
S105	Jane	Brown	S103	IAI	58
S105	Jane	Brown	S104	PR1	68
S105	Jane	Brown	S104	IAI	65
S105	Jane	Brown	S106	PR2	43
S105	Jane	Brown	S107	PR1	76
S105	Jane	Brown	S107	PR2	60
S105	Jane	Brown	S107	IAI	35

```
SELECT *  
FROM student, grade  
WHERE student.id = grade.id
```

id	first	last	id	code	mark
S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58
S104	Mary	Jones	S104	PR1	68
S104	Mary	Jones	S104	IAI	65
S106	Mark	Jones	S106	PR2	43
S107	John	Brown	S107	PR1	76
S107	John	Brown	S107	PR2	60
S107	John	Brown	S107	IAI	35

After you have a single table!

And you know how to deal with those...

```
SELECT *
FROM student, grade
WHERE (student.id = grade.id)
```

Graded Students

id	first	last	id	code	mark
S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58
S104	Mary	Jones	S104	PR1	68
S104	Mary	Jones	S104	IAI	65
S106	Mark	Jones	S106	PR2	43
S107	John	Brown	S107	PR1	76
S107	John	Brown	S107	PR2	60
S107	John	Brown	S107	IAI	35



After you have a single table!

And you know how to deal with those...

```
SELECT *  
FROM student, grade  
WHERE (student.id = grade.id)
```

```
SELECT first, last, code, mark  
FROM student, grade  
WHERE (student.id = grade.id)
```

Graded Students

id	first	last	id	code	mark
S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58
S104	Mary	Jones	S104	PR1	68
S104	Mary	Jones	S104	IAI	65
S106	Mark	Jones	S106	PR2	43
S107	John	Brown	S107	PR1	76
S107	John	Brown	S107	PR2	60
S107	John	Brown	S107	IAI	35



first	last	code	mark
John	Smith	DBS	72
John	Smith	IAI	58
Mary	Jones	PR1	68
Mary	Jones	IAI	65
Mark	Jones	PR2	43
John	Brown	PR1	76
John	Brown	PR2	60
John	Brown	IAI	35

After you have a single table!

And you know how to deal with those...

```
SELECT *  
FROM student, grade  
WHERE (student.id = grade.id)
```

```
SELECT first, last, code, mark  
FROM student, grade  
WHERE (student.id = grade.id)
```

```
SELECT first, last  
FROM student, grade  
WHERE (student.id = grade.id)  
AND code = 'PR1'  
AND mark > 70
```

Graded Students

id	first	last	id	code	mark
S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58
S104	Mary	Jones	S104	PR1	68
S104	Mary	Jones	S104	IAI	65
S106	Mark	Jones	S106	PR2	43
S107	John	Brown	S107	PR1	76
S107	John	Brown	S107	PR2	60
S107	John	Brown	S107	IAI	35

first	last	code	mark
John	Smith	DBS	72
John	Smith	IAI	58
Mary	Jones	PR1	68
Mary	Jones	IAI	65
Mark	Jones	PR2	43
John	Brown	PR1	76
John	Brown	PR2	60
John	Brown	IAI	35

first	last
John	Brown

After we exhaustively join first two tables, let's exhaustively joint the third...

```
SELECT *  
FROM student, grade, course
```

id	first	last	id	code	mark
S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58
S103	John	Smith	S104	PR1	68
S103	John	Smith	S104	IAI	65
S103	John	Smith	S106	PR2	43
S103	John	Smith	S107	PR1	76
S103	John	Smith	S107	PR2	60
S103	John	Smith	S107	IAI	35
S104	Mary	Jones	S103	DBS	72
S104	Mary	Jones	S103	IAI	58
S104	Mary	Jones	S104	PR1	68
S104	Mary	Jones	S104	IAI	65
S104	Mary	Jones	S106	PR2	43
S104	Mary	Jones	S107	PR1	76
S104	Mary	Jones	S107	PR2	60
S104	Mary	Jones	S107	IAI	35



After we exhaustively join first two tables exhaustively joint the third...

```
SELECT *  
FROM student, grade, course
```

id	first	last	id	code	mark
S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58
S103	John	Smith	S104	PR1	68
S103	John	Smith	S104	IAI	65
S103	John	Smith	S106	PR2	43
S103	John	Smith	S107	PR1	76
S103	John	Smith	S107	PR2	60
S103	John	Smith	S107	IAI	35
S104	Mary	Jones	S103	DBS	72
S104	Mary	Jones	S103	IAI	58
S104	Mary	Jones	S104	PR1	68
S104	Mary	Jones	S104	IAI	65
S104	Mary	Jones	S106	PR2	43
S104	Mary	Jones	S107	PR1	76
S104	Mary	Jones	S107	PR2	60
S104	Mary	Jones	S107	IAI	35

id	first	last	id	code	mark	code	title
S103	John	Smith	S103	DBS	72	DBS	Database Systems
						PR1	Programming 1
						PR2	Programming 2
						IAI	Intro to AI
S103	John	Smith	S103	IAI	58	DBS	Database Systems
						PR1	Programming 1
						PR2	Programming 2
						IAI	Intro to AI
S103	John	Smith	S104	PR1	68	DBS	Database Systems
						PR1	Programming 1
						PR2	Programming 2
						IAI	Intro to AI
S103	John	Smith	S104	IAI	65	DBS	Database Systems
						PR1	Programming 1
						PR2	Programming 2
						IAI	Intro to AI
S103	John	Smith	S106	PR2	43	DBS	Database Systems
						PR1	Programming 1
						PR2	Programming 2
						IAI	Intro to AI
S103	John	Smith	S107	PR1	76	DBS	Database Systems
						PR1	Programming 1
						PR2	Programming 2
						IAI	Intro to AI

course

code	title
DBS	Database Systems
PR1	Programming 1
PR2	Programming 2
IAI	Intro to AI

SQL: JOINING TABLES

(the traditional way)

grade

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68
S104	IAI	65
S106	PR2	43
S107	PR1	76
S107	PR2	60
S107	IAI	35

student

id	first	last
S103	John	Smith
S104	Mary	Jones
S105	Jane	Brown
S106	Mark	Jones
S107	John	Brown

course

code	title
DBS	Database Systems
PR1	Programming 1
PR2	Programming 2
IAI	Intro to AI

What about three tables?

Same deal, facts "new" entity.

Joining student: S103, John, Smith

With grade: S103, DBS, 72

With course: DBS, Database Systems

Creates a new, "**detailed graded student**" entity.

id	first	last	id	code	mark	code	title
S103	John	Smith	S103	DBS	72	DBS	Database Systems
S103	John	Smith	S103	DBS	72	PR1	Programming 1
S103	John	Smith	S103	DBS	72	PR2	Programming 2
S103	John	Smith	S103	DBS	72	IAI	Intro to AI
S103	John	Smith	S103	IAI	58	DBS	Database Systems
S103	John	Smith	S103	IAI	58	PR1	Programming 1
S103	John	Smith	S103	IAI	58	PR2	Programming 2
S103	John	Smith	S103	IAI	58	IAI	Intro to AI
S103	John	Smith	S104	PR1	68	DBS	Database Systems
S103	John	Smith	S104	PR1	68	PR1	Programming 1
S103	John	Smith	S104	PR1	68	PR2	Programming 2
S103	John	Smith	S104	PR1	68	IAI	Intro to AI
S103	John	Smith	S104	IAI	65	DBS	Database Systems
S103	John	Smith	S104	IAI	65	PR1	Programming 1
S103	John	Smith	S104	IAI	65	PR2	Programming 2
S103	John	Smith	S104	IAI	65	IAI	Intro to AI
S103	John	Smith	S106	PR2	43	DBS	Database Systems
S103	John	Smith	S106	PR2	43	PR1	Programming 1
S103	John	Smith	S106	PR2	43	PR2	Programming 2
S103	John	Smith	S106	PR2	43	IAI	Intro to AI
S103	John	Smith	S107	PR1	76	DBS	Database Systems

Again, we must filter to retain only rows that are **true**

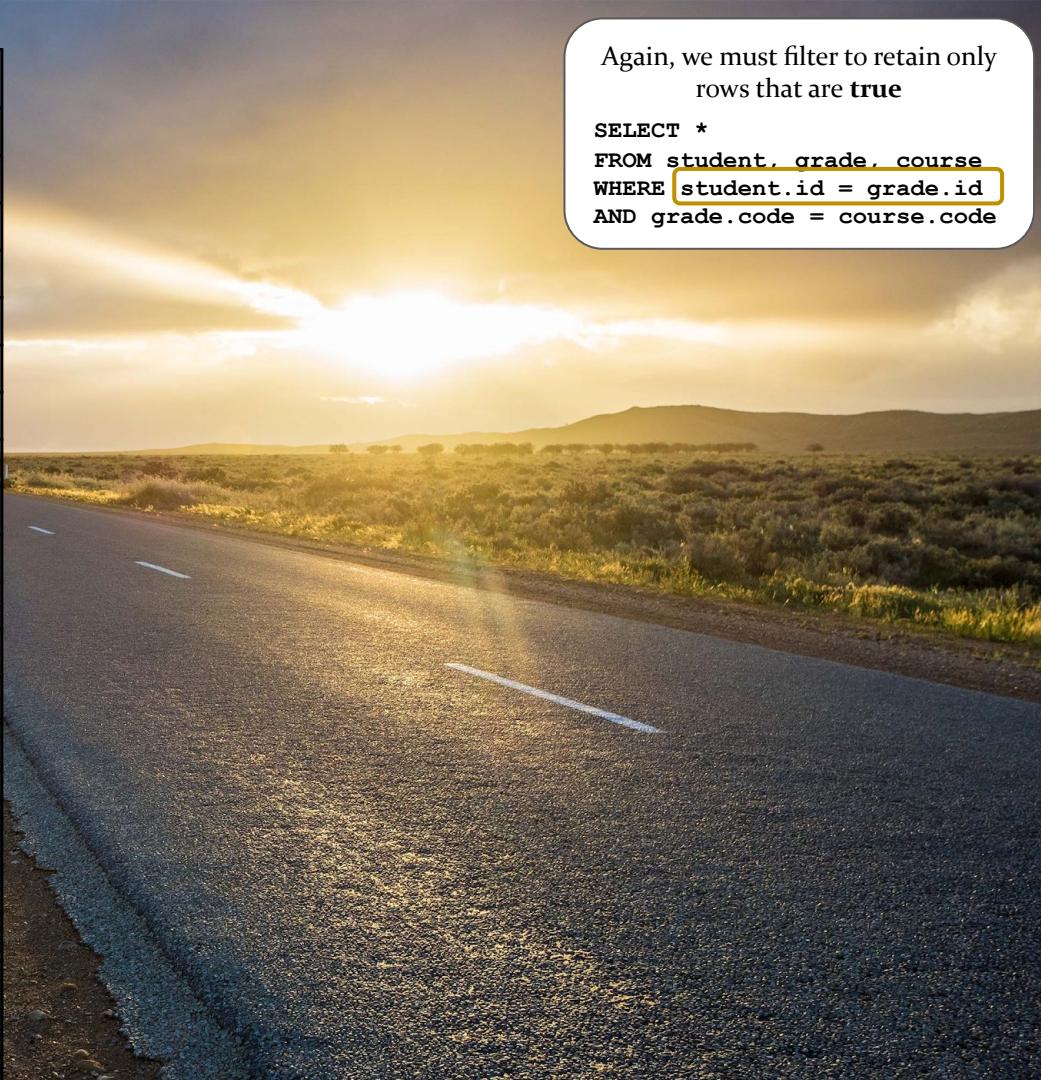
```
SELECT *
FROM student, grade, course
WHERE student.id = grade.id
AND grade.code = course.code
```



id	first	last	id	code	mark	code	title
S103	John	Smith	S103	DBS	72	DBS	Database Systems
S103	John	Smith	S103	DBS	72	PR1	Programming 1
S103	John	Smith	S103	DBS	72	PR2	Programming 2
S103	John	Smith	S103	DBS	72	IAI	Intro to AI
S103	John	Smith	S103	IAI	58	DBS	Database Systems
S103	John	Smith	S103	IAI	58	PR1	Programming 1
S103	John	Smith	S103	IAI	58	PR2	Programming 2
S103	John	Smith	S103	IAI	58	IAI	Intro to AI
S103	John	Smith	S104	PR1	68	DBS	Database Systems
S103	John	Smith	S104	PR1	68	PR1	Programming 1
S103	John	Smith	S104	PR1	68	PR2	Programming 2
S103	John	Smith	S104	PR1	68	IAI	Intro to AI
S103	John	Smith	S104	IAI	65	DBS	Database Systems
S103	John	Smith	S104	IAI	65	PR1	Programming 1
S103	John	Smith	S104	IAI	65	PR2	Programming 2
S103	John	Smith	S104	IAI	65	IAI	Intro to AI
S103	John	Smith	S106	PR2	43	DBS	Database Systems
S103	John	Smith	S106	PR2	43	PR1	Programming 1
S103	John	Smith	S106	PR2	43	PR2	Programming 2
S103	John	Smith	S106	PR2	43	IAI	Intro to AI
S103	John	Smith	S107	PR1	76	DBS	Database Systems

Again, we must filter to retain only rows that are **true**

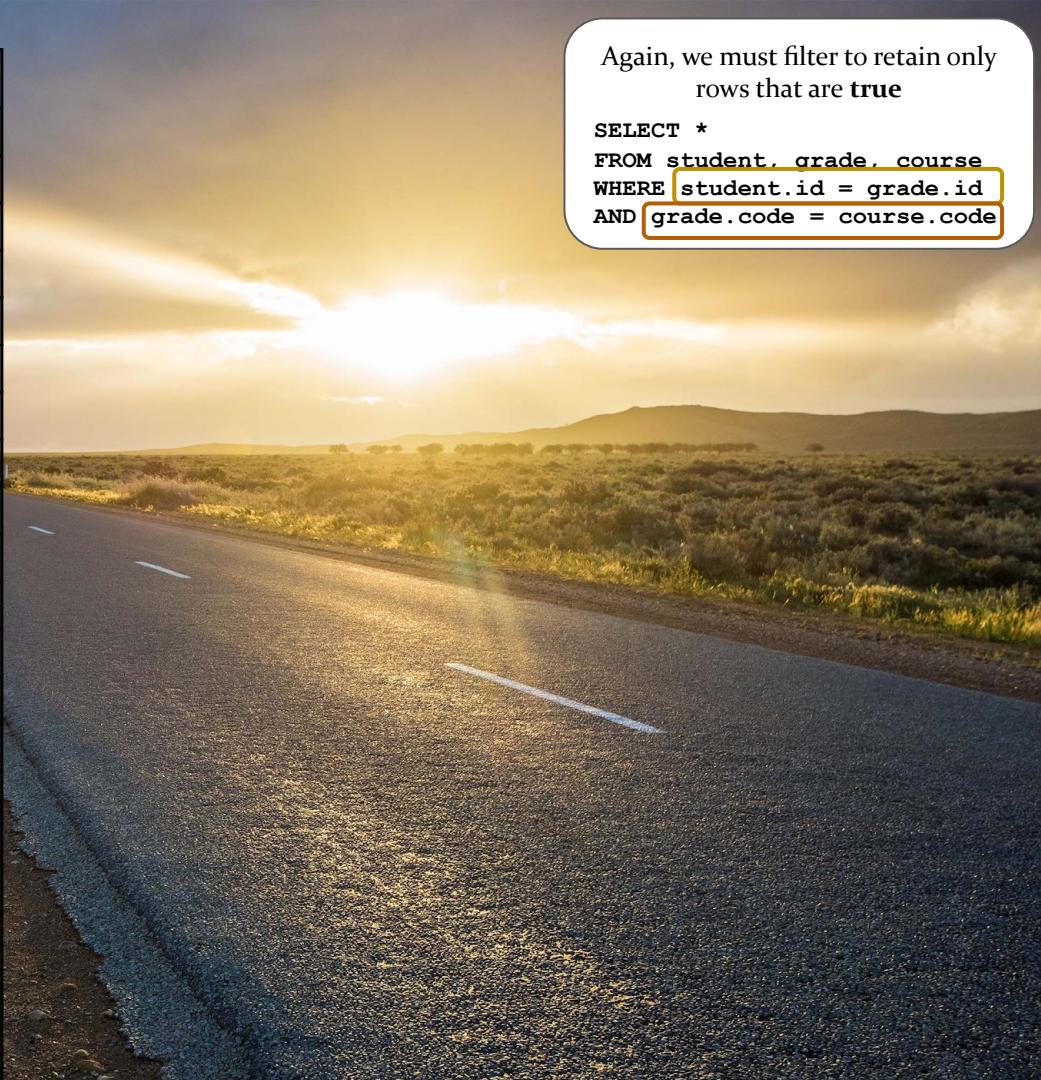
```
SELECT *
FROM student, grade, course
WHERE student.id = grade.id
AND grade.code = course.code
```



id	first	last	id	code	mark	code	title
S103	John	Smith	S103	DBS	72	DBS	Database Systems
S103	John	Smith	S103	DBS	72	PR1	Programming 1
S103	John	Smith	S103	DBS	72	PR2	Programming 2
S103	John	Smith	S103	DBS	72	IAI	Intro to AI
S103	John	Smith	S103	IAI	58	DBS	Database Systems
S103	John	Smith	S103	IAI	58	PR1	Programming 1
S103	John	Smith	S103	IAI	58	PR2	Programming 2
S103	John	Smith	S103	IAI	58	IAI	Intro to AI
S103	John	Smith	S104	PR1	68	DBS	Database Systems
S103	John	Smith	S104	PR1	68	PR1	Programming 1
S103	John	Smith	S104	PR1	68	PR2	Programming 2
S103	John	Smith	S104	PR1	68	IAI	Intro to AI
S103	John	Smith	S104	IAI	65	DBS	Database Systems
S103	John	Smith	S104	IAI	65	PR1	Programming 1
S103	John	Smith	S104	IAI	65	PR2	Programming 2
S103	John	Smith	S104	IAI	65	IAI	Intro to AI
S103	John	Smith	S106	PR2	43	DBS	Database Systems
S103	John	Smith	S106	PR2	43	PR1	Programming 1
S103	John	Smith	S106	PR2	43	PR2	Programming 2
S103	John	Smith	S106	PR2	43	IAI	Intro to AI
S103	John	Smith	S107	PR1	76	DBS	Database Systems

Again, we must filter to retain only rows that are **true**

```
SELECT *
FROM student, grade, course
WHERE student.id = grade.id
AND grade.code = course.code
```



id	first	last	id	code	mark	code	title
S103	John	Smith	S103	DBS	72	DBS	Database Systems
S103	John	Smith	S103	DBS	72	PR1	Programming 1
S103	John	Smith	S103	DBS	72	PR2	Programming 2
S103	John	Smith	S103	DBS	72	IAI	Intro to AI
S103	John	Smith	S103	IAI	58	DBS	Database Systems
S103	John	Smith	S103	IAI	58	PR1	Programming 1
S103	John	Smith	S103	IAI	58	PR2	Programming 2
S103	John	Smith	S103	IAI	58	IAI	Intro to AI
S103	John	Smith	S104	PR1	68	DBS	Database Systems
S103	John	Smith	S104	PR1	68	PR1	Programming 1
S103	John	Smith	S104	PR1	68	PR2	Programming 2
S103	John	Smith	S104	PR1	68	IAI	Intro to AI
S103	John	Smith	S104	IAI	65	DBS	Database Systems
S103	John	Smith	S104	IAI	65	PR1	Programming 1
S103	John	Smith	S104	IAI	65	PR2	Programming 2
S103	John	Smith	S104	IAI	65	IAI	Intro to AI
S103	John	Smith	S106	PR2	43	DBS	Database Systems
S103	John	Smith	S106	PR2	43	PR1	Programming 1
S103	John	Smith	S106	PR2	43	PR2	Programming 2
S103	John	Smith	S106	PR2	43	IAI	Intro to AI
S103	John	Smith	S107	PR1	76	DBS	Database Systems

✓ ✗

Again, we must filter to retain only rows that are **true**

```
SELECT *
FROM student, grade, course
WHERE student.id = grade.id
AND grade.code = course.code
```

id	first	last	id	code	mark	code	title
S103	John	Smith	S103	DBS	72	DBS	Database Systems
S103	John	Smith	S103	DBS	72	PR1	Programming 1
S103	John	Smith	S103	DBS	72	PR2	Programming 2
S103	John	Smith	S103	DBS	72	IAI	Intro to AI
S103	John	Smith	S103	IAI	58	DBS	Database Systems
S103	John	Smith	S103	IAI	58	PR1	Programming 1
S103	John	Smith	S103	IAI	58	PR2	Programming 2
S103	John	Smith	S103	IAI	58	IAI	Intro to AI
S103	John	Smith	S104	PR1	68	DBS	Database Systems
S103	John	Smith	S104	PR1	68	PR1	Programming 1
S103	John	Smith	S104	PR1	68	PR2	Programming 2
S103	John	Smith	S104	PR1	68	IAI	Intro to AI
S103	John	Smith	S104	IAI	65	DBS	Database Systems
S103	John	Smith	S104	IAI	65	PR1	Programming 1
S103	John	Smith	S104	IAI	65	PR2	Programming 2
S103	John	Smith	S104	IAI	65	IAI	Intro to AI
S103	John	Smith	S106	PR2	43	DBS	Database Systems
S103	John	Smith	S106	PR2	43	PR1	Programming 1
S103	John	Smith	S106	PR2	43	PR2	Programming 2
S103	John	Smith	S106	PR2	43	IAI	Intro to AI
S103	John	Smith	S107	PR1	76	DBS	Database Systems

Again, we must filter to retain only rows that are **true**

```
SELECT *
FROM student, grade, course
WHERE student.id = grade.id
AND grade.code = course.code
```

Again, we must filter to retain only rows that are **true**

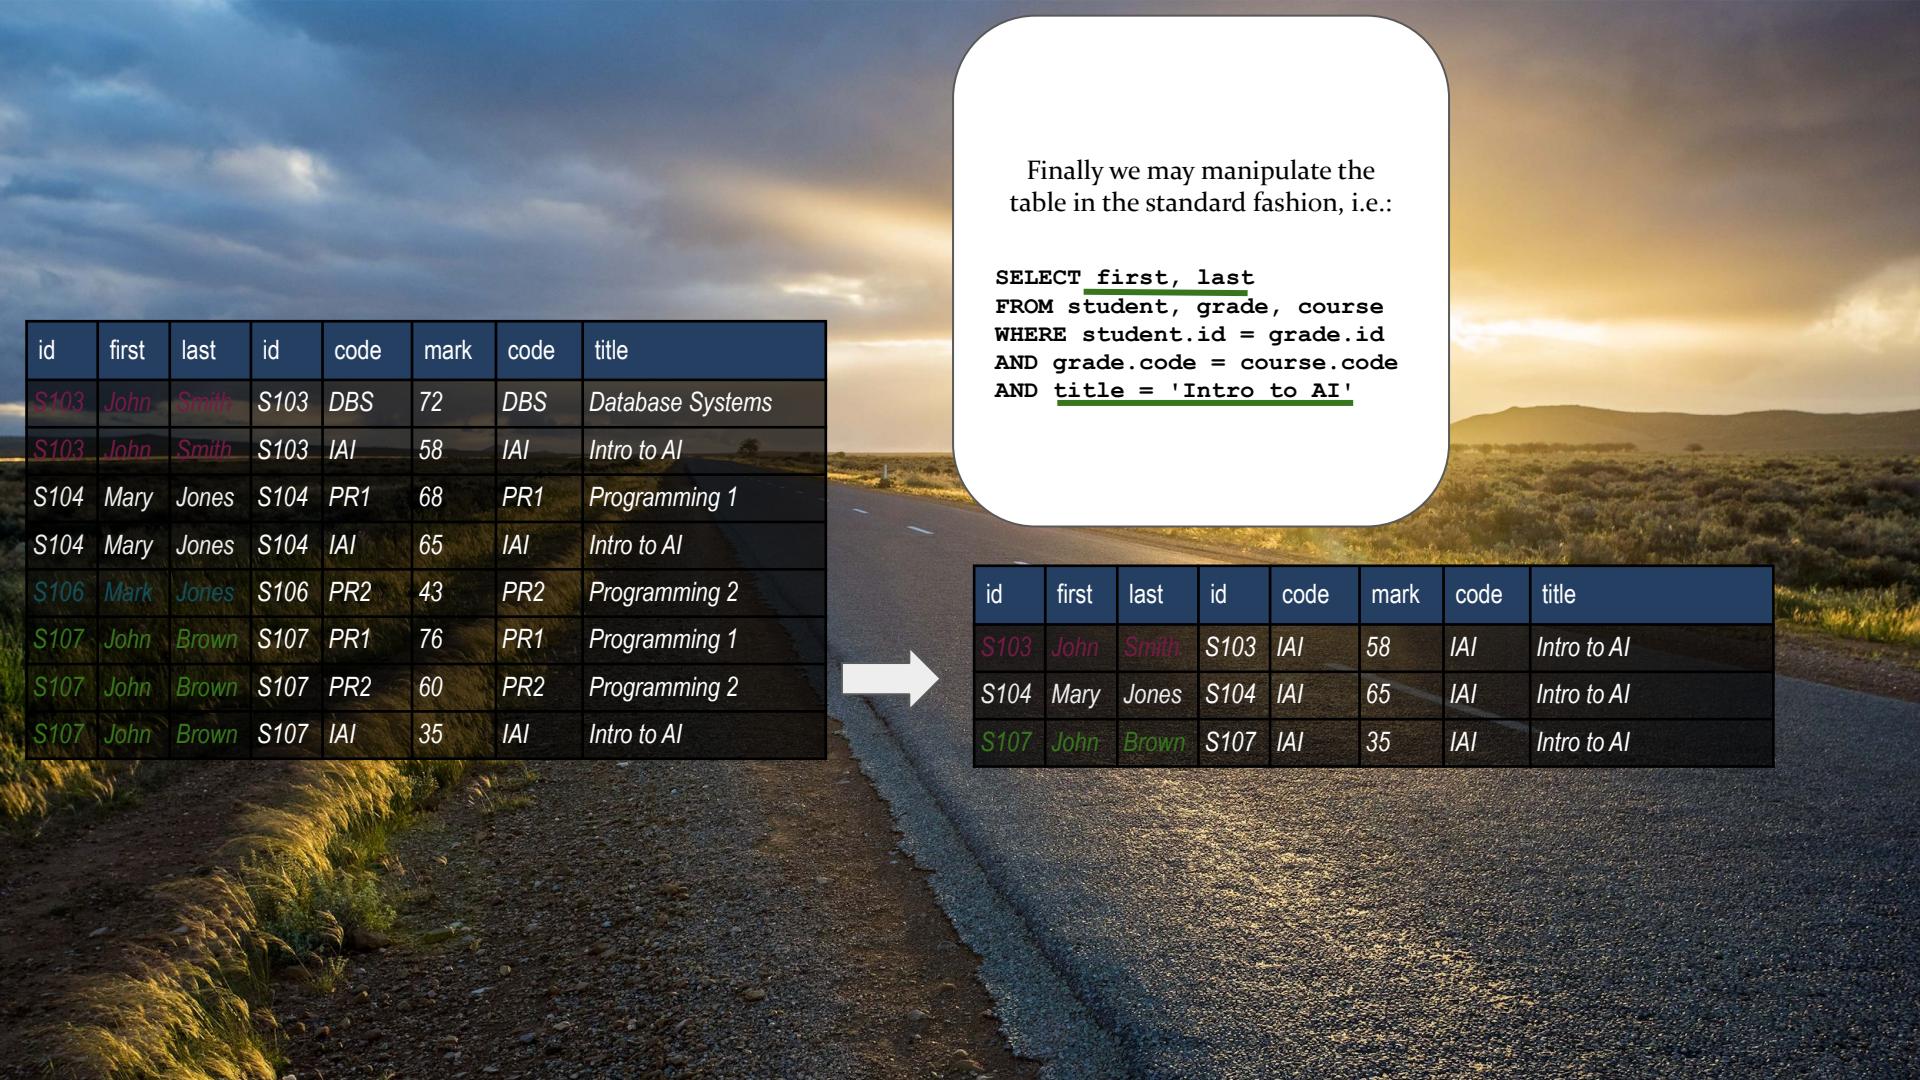
```
SELECT *
FROM student, grade, course
WHERE student.id = grade.id
AND grade.code = course.code
```

id	first	last	id	code	mark	code	title
S103	John	Smith	S103	DBS	72	DBS	Database Systems
S103	John	Smith	S103	DBS	72	PR1	Programming 1
S103	John	Smith	S103	DBS	72	PR2	Programming 2
S103	John	Smith	S103	DBS	72	IAI	Intro to AI
S103	John	Smith	S103	IAI	58	DBS	Database Systems
S103	John	Smith	S103	IAI	58	PR1	Programming 1
S103	John	Smith	S103	IAI	58	PR2	Programming 2
S103	John	Smith	S103	IAI	58	IAI	Intro to AI
S103	John	Smith	S104	PR1	68	DBS	Database Systems
S103	John	Smith	S104	PR1	68	PR1	Programming 1
S103	John	Smith	S104	PR1	68	PR2	Programming 2
S103	John	Smith	S104	PR1	68	IAI	Intro to AI
S103	John	Smith	S104	IAI	65	DBS	Database Systems
S103	John	Smith	S104	IAI	65	PR1	Programming 1
S103	John	Smith	S104	IAI	65	PR2	Programming 2
S103	John	Smith	S104	IAI	65	IAI	Intro to AI
S103	John	Smith	S106	PR2	43	DBS	Database Systems
S103	John	Smith	S106	PR2	43	PR1	Programming 1
S103	John	Smith	S106	PR2	43	PR2	Programming 2
S103	John	Smith	S106	PR2	43	IAI	Intro to AI
S103	John	Smith	S107	PR1	76	DBS	Database Systems



id	first	last	id	code	mark	code	title
S103	John	Smith	S103	DBS	72	DBS	Database Systems
S103	John	Smith	S103	IAI	58	IAI	Intro to AI
S104	Mary	Jones	S104	PR1	68	PR1	Programming 1
S104	Mary	Jones	S104	IAI	65	IAI	Intro to AI
S106	Mark	Jones	S106	PR2	43	PR2	Programming 2
S107	John	Brown	S107	PR1	76	PR1	Programming 1
S107	John	Brown	S107	PR2	60	PR2	Programming 2
S107	John	Brown	S107	IAI	35	IAI	Intro to AI





Finally we may manipulate the table in the standard fashion, i.e.:

```
SELECT first, last
FROM student, grade, course
WHERE student.id = grade.id
AND grade.code = course.code
AND title = 'Intro to AI'
```

id	first	last	id	code	mark	code	title
S103	John	Smith	S103	DBS	72	DBS	Database Systems
S103	John	Smith	S103	IAI	58	IAI	Intro to AI
S104	Mary	Jones	S104	PR1	68	PR1	Programming 1
S104	Mary	Jones	S104	IAI	65	IAI	Intro to AI
S106	Mark	Jones	S106	PR2	43	PR2	Programming 2
S107	John	Brown	S107	PR1	76	PR1	Programming 1
S107	John	Brown	S107	PR2	60	PR2	Programming 2
S107	John	Brown	S107	IAI	35	IAI	Intro to AI



id	first	last	id	code	mark	code	title
S103	John	Smith	S103	IAI	58	IAI	Intro to AI
S104	Mary	Jones	S104	IAI	65	IAI	Intro to AI
S107	John	Brown	S107	IAI	35	IAI	Intro to AI

Some observations...

grade

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68
S104	IAI	65
S106	PR2	43
S107	PR1	76
S107	PR2	60
S107	IAI	35

students

id	first	last	age
S103	John	Smith	22
S104	Mary	Jones	27
S105	Jane	Brown	21
S106	Mark	Jones	26
S107	John	Brown	30

course

code	title
DBS	Database Systems
PR1	Programming 1
PR2	Programming 2
IAI	Intro to AI

Joining tables gives us:

New set of entities

(row: "Graded student", new set of jointly true facts)



Some observations...

grade

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68
S104	IAI	65
S106	PR2	43
S107	PR1	76
S107	PR2	60
S107	IAI	35

students

id	first	last	age
S103	John	Smith	22
S104	Mary	Jones	27
S105	Jane	Brown	21
S106	Mark	Jones	26
S107	John	Brown	30

course

code	title
DBS	Database Systems
PR1	Programming 1
PR2	Programming 2
IAI	Intro to AI

Joining tables gives us:

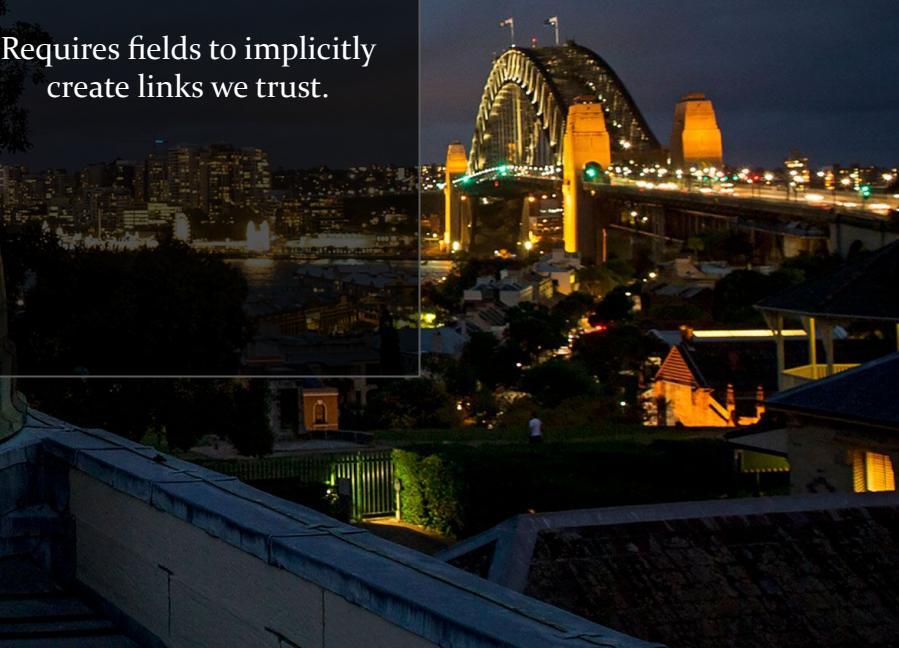
New set of entities

(row: "Graded student", new set of jointly true facts)

Does the joint entity you're making make sense?

Will the each **rows facts** be **jointly true?**

Requires fields to implicitly create links we trust.



Some observations...

grade

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68
S104	IAI	65
S106	PR2	43
S107	PR1	76
S107	PR2	60
S107	IAI	35

students

id	first	last	age
S103	John	Smith	22
S104	Mary	Jones	27
S105	Jane	Brown	21
S106	Mark	Jones	26
S107	John	Brown	30

course

code	title
DBS	Database Systems
PR1	Programming 1
PR2	Programming 2
IAI	Intro to AI

Joining tables gives us:

New set of entities

(row: "Graded student", new set of jointly true facts)

Does the joint entity you're making make sense?
Will the each **rows** facts be **jointly true**?

Requires fields to implicitly create links we trust.

Links considered at **design-time**: unambiguous

(e.g. Bob's ID in hospital & chef table
e.g. transaction_id in basket and line_items table)

Other links (e.g. linking by date): Opportunities! But think carefully...



Understanding when to link: Documentation

Documentation describes
fields that link data and
tell us **the relationship**
(how to interpret the result after linking)

Recall:

Candidate/Primary keys:
Uniquely identify a entity
occurrence (row)

Foreign keys:
Record an identifier (link) to
a candidate key in another
table

Understanding when to link: Documentation

Documentation describes
fields that link data and
tell us **the relationship**
(how to interpret the result after linking)

How we **interpret** the resulting **entity instances** is guided by the relationship information (documented = enforced) (non-documented = at our own risk)

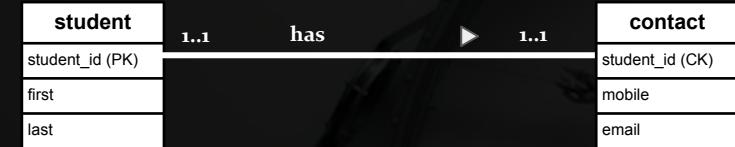
Recall:

Candidate/Primary keys:
Uniquely identify a entity occurrence (row)

Foreign keys:
Record an identifier (link) to a candidate key in another table

One-to-one

Relationship is always implicit,
and is realised due to the JOIN
and the constraints (enforced by
design or opportunistic) on to the
linking fields.



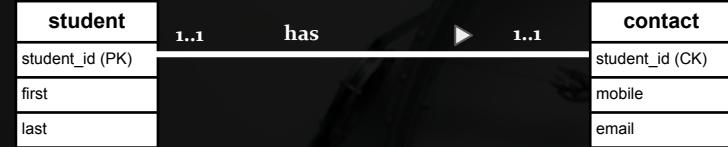
id	first	last
1	Gav	Smith
2	Bob	Brown

id	mobile	email
1	123	g@s
2	586	b@b

One-to-one

Relationship is always implicit, and is realised due to the JOIN and the constraints (enforced by design or opportunistic) on to the linking fields.

Consider the join of the tables in this example. All rows from **student** are joined with all rows from **contact**.



id	first	last
1	Gav	Smith
2	Bob	Brown

id	mobile	email
1	123	g@s
2	586	b@b

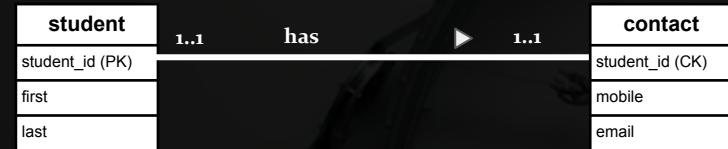
id	first	last	id	mobile	email
1	Gav	Smith	1	123	g@s
1	Gav	Smith	2	586	b@b
2	Bob	Brown	1	123	g@s
2	Bob	Brown	2	586	b@b

One-to-one

Relationship is always implicit, and is realised due to the JOIN and the constraints (enforced by design or opportunistic) on to the linking fields.

Consider the join of the tables in this example. All rows from **student** are joined with all rows from **contact**.

Due to the PK and CK constraints on the linking fields, only one row per student possible in the output table.



id	first	last
1	Gav	Smith
2	Bob	Brown

id	mobile	email
1	123	g@s
2	586	b@b

id	first	last	id	mobile	email
1	Gav	Smith	1	123	g@s
1	Gav	Smith	2	586	b@b
2	Bob	Brown	1	123	g@s
2	Bob	Brown	2	586	b@b

id	first	last	mobile	email
1	Gav	Smith	123	g@s
2	Bob	Brown	486	b@b

One-to-one

(enforced by cust_id being primary/candidate keys in each table)

Relationship is always implicit, and is realised due to the JOIN and the constraints (enforced by design or opportunistic) on to the linking fields.

Consider the join of the tables in this example. All rows from **student** are joined with all rows from **contact**.

Due to the PK and CK constraints on the linking fields, only one row per student possible in the output table.



Interpretation is simple. Equivalently:

- Each **contact** record is **extended** with **student** information.
- Each **student** record is **extended** with **contact** information.

id	first	last	mobile	email
1	Gav	Smith	123	g@s
2	Bob	Brown	486	b@b

Most Logical (maybe) Interpretation

(but not the easiest to use to "think" in SQL as it is a conceptual nested structure - not the basis of SQL!)

Each customer has a list of transactions.

Could view this as a customer "entity" with multiple transactions.

Mark

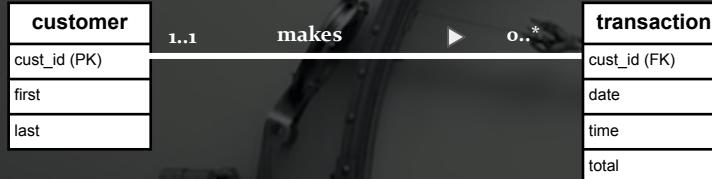
→	2017-08-09	08:05:32	£12.00
→	2017-08-10	13:00:00	£6.50

Jane

→	2017-08-10	11:00:01	£3.00
---	------------	----------	-------

One-to-many

(enforced by cust_id being primary key in the "one end" of the table and a foreign key in the "many end" of the relationship)



Entity type: customer

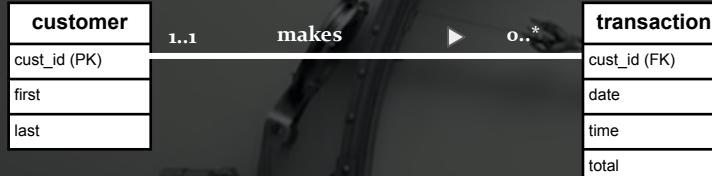
cust_id	first	last
1	Mark	Jones
2	Victoria	Smith
3	Jane	Doe

Entity type: transaction

cust_id	date	time	total
1	2017-08-09	08:05:32	£12.00
1	2017-08-10	13:00:00	£6.50
3	2017-08-10	11:00:01	£3.00

One-to-many

(enforced by `cust_id` being primary key in the "one end" of the table and a foreign key in the "many end" of the relationship)



Most Logical (maybe) Interpretation

(but not the easiest to use to "think" in SQL as it is a conceptual nested structure - not the basis of SQL!)

Each customer has a list of transactions.

Could view this as a customer "entity" with multiple transactions.

Mark

→	2017-08-09	08:05:32	£12.00
→	2017-08-10	13:00:00	£6.50

Jane

→	2017-08-10	11:00:01	£3.00
---	------------	----------	-------

Entity type: customer			Entity type: transaction			
cust_id	first	last	cust_id	date	time	total
1	Mark	Jones	1	2017-08-09	08:05:32	£12.00
2	Victoria	Smith	1	2017-08-10	13:00:00	£6.50
3	Jane	Doe	3	2017-08-10	11:00:01	£3.00

Relationship is always implicit, and is realised due to the JOIN and the constraints (enforced by design or opportunistic) on to the linking fields.

Enforced by

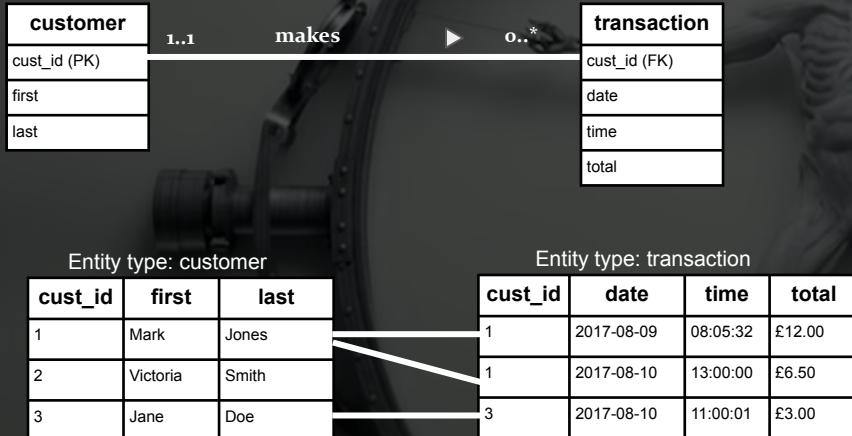
→ **`cust_id` being primary key in the "one end" of the table**

→ **a foreign key in the "many end" of the relationship**
(records the `cust_id` of the transaction)

→ When all rows from each table are compared, only true rows are those where the customer id's match.

I.e. where the transaction record has been extended with the right customer information.

One-to-many



Potentially easier
if we stay "thinking
in tables"

Alternate interpretation
View as a transaction
entity augmented with
customer information.

cust_id	first	last	date	time	total
1	Mark	Jones	2017-08-09	08:05:32	£12.00
1	Mark	Jones	2017-08-10	13:00:00	£6.50
3	Jane	Doe	2017-08-10	11:00:01	£3.00

One-to-many

Remember:

The result of joining two tables is a new table.

**Always think:
What is the entity type represented?**

Relationships (particularly if they are enforced) guide our interpretation.

Each customer has a list of transactions.

Could view this as a customer "entity" with multiple transactions.

Mark	→	2017-08-09	08:05:32	£12.00
	→	2017-08-10	13:00:00	£6.50
Jane	→	2017-08-10	11:00:01	£3.00

A
V
E
C

Identical structure information. But we have a table. **Keeping things simple for SQL.**

cust_id	first	last	date	time	total
1	Mark	Jones	2017-08-09	08:05:32	£12.00
			2017-08-10	13:00:00	£6.50
3	Jane	Doe	2017-08-10	11:00:01	£3.00

Many-to-many

Entity type: basket

b_id	time	loyalty_id	total
b1	11am	7896	£3.70
b2	1pm	9934	£3.25
b3	2pm	3900	£0.07

Entity type: items

i_id	desc.	cost
i1	Doritos	£3.00
i2	Chocolate bar	£0.70
i3	Sandwich	£3.25



Should not exist in a well designed relational database.

Interpretation

A basket entity (row) would need to store many items. A row must, however, only be atomic entities.

Entity type: basket

b_id	time	loyalty_id	total
b1	11am	7896	£3.70
b2	1pm	9934	£3.25
b3	2pm	3900	£0.07

Entity type: items

i_id	desc.	cost
i1	Doritos	£3.00
i2	Chocolate bar	£0.70
i3	Sandwich	£3.25



Many-to-many

Interpretation

Should not exist in a well designed relational database.

If it did, the items table is no longer "items" as we'd conceptualize it (one item is one entity occurrence).

Entity type: basket

b_id	time	loyalty_id	total
b1	11am	7896	£3.70
b2	1pm	9934	£3.25
b3	2pm	3900	£0.07

Entity type: items

i_id	desc.	cost
i1	Doritos	£3.00
i2	Chocolate bar	£0.70
i3	Sandwich	£3.25

basket

time
loyalty_id
total

items

desc.
cost

o..*



1..*

items

basket
basket_id
time
loyalty_id
total

line_items
basket_id
item_id

*

o..*



1..*

listed in

items
item_id
desc.
cost

Same information, corrected to being actual "tables". Becomes three tables and two one-to-many relationships.

AGGREGATE FUNCTIONS

grade

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68
S104	IAI	65
S106	PR2	43
S107	PR1	76
S107	PR2	60
S107	IAI	35

OK, so joins are fun...
But let's step back...

And look at operations on
a single (or post join)
table.

Question: What is the
maximum mark per
course?

Answer in SQL:

```
SELECT code, MAX(mark)
FROM grade
GROUP BY code
```

AGGREGATE FUNCTIONS

grade

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68
S104	IAI	65
S106	PR2	43
S107	PR1	76
S107	PR2	60
S107	IAI	35

OK, so joins are fun...
But let's step back...

And look at operations on
a single (or post join)
table.

Question: What is the
maximum mark per
course?

Answer in SQL:

```
SELECT code, MAX(mark)
FROM grade
GROUP BY code
```

Step 1: Group (put) the
rows into buckets, one
bucket per distinct module
code value

AGGREGATE FUNCTIONS

grade

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68
S104	IAI	65
S106	PR2	43
S107	PR1	76
S107	PR2	60
S107	IAI	35

OK, so joins are fun...
But let's step back...

Question: What is the maximum mark per course?

And look at operations on a single (or post join) table.

Answer in SQL:

```
SELECT code, MAX(mark)  
FROM grade  
GROUP BY code
```

Step 2: Per group (bucket) compute the aggregate function (in this case max) over all values in the bucket

Step 1: Group (put) the rows into buckets, one bucket per distinct module code value

AGGREGATE FUNCTIONS

grade

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68
S104	IAI	65
S106	PR2	43
S107	PR1	76
S107	PR2	60
S107	IAI	35

OK, so joins are fun...
But let's step back...

Question: What is the maximum mark per course?

And look at operations on a single (or post join) table.

Answer in SQL:

```
SELECT code, MAX(mark)  
FROM grade  
GROUP BY code
```

Step 3: Return the bucket label and the aggregate values.

NOTE: Bucket labels need not be included. Fields not part of a bucket label can't be.

WHY? Ambiguous, what value of the many should have been picked per row?

Step 2: Per group (bucket) compute the aggregate function (in this case max) over all values in the bucket

Step 1: Group (put) the rows into buckets, one bucket per distinct module code value

AGGREGATE FUNCTIONS

grade

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68
S104	IAI	65
S106	PR2	43
S107	PR1	76
S107	PR2	60
S107	IAI	35

What about if we want just
one big bucket?

(group over everything)
(max over all rows)

Question: What is the
maximum mark?

Answer in SQL:

```
SELECT MAX(mark)  
FROM grade
```

AGGREGATE FUNCTIONS

grade

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68
S104	IAI	65
S106	PR2	43
S107	PR1	76
S107	PR2	60
S107	IAI	35

What about if we want just one big bucket?

(group over everything)
(max over all rows)

Question: What is the maximum mark?

Answer in SQL:

```
SELECT MAX(mark)  
FROM grade
```

Step 0: Don't specify bucket labels.

AGGREGATE FUNCTIONS

grade

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68
S104	IAI	65
S106	PR2	43
S107	PR1	76
S107	PR2	60
S107	IAI	35

What about if we want just one big bucket?

(group over everything)
(max over all rows)

Question: What is the maximum mark?

Answer in SQL:

```
SELECT  
FROM grade
```

MAX (mark)

Step 2: The aggregate function (in this case max) is computed over all values.

Step 0: Don't specify bucket labels.

AGGREGATE FUNCTIONS

grade

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68
S104	IAI	65
S106	PR2	43
S107	PR1	76
S107	PR2	60
S107	IAI	35

What about if we want just one big bucket?

(group over everything)
(max over all rows)

Question: What is the maximum mark?

Answer in SQL:

```
SELECT MAX(mark)  
FROM grade
```

Step 3: Return the aggregate values

NOTE: No bucket labels, no non-aggregate terms.

WHY? Ambiguous, what value of the many should have been picked per row?

Step 2: The aggregate function (in this case max) is computed over all values.

Step 1: Don't specify bucket labels.

AGGREGATE FUNCTIONS

grade

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68
S104	IAI	65
S106	PR2	43
S107	PR1	76
S107	PR2	60
S107	IAI	35

There are many aggregate functions.

You can write your own
(in Python even!)

Standard aggregate functions

count avg variance
max sum median
min stdev

AGGREGATE FUNCTIONS

grade

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68
S104	IAI	65
S106	PR2	43
S107	PR1	76
S107	PR2	60
S107	IAI	35

All operate under the same principal.

Work on the bucket of data given by the group by.

More than one aggregate can be run over the buckets at once.

```
SELECT code,  
       MIN(mark),  
       MAX(mark),  
       AVG(mark),  
FROM grade  
GROUP BY code
```

id	code	mark
S103	DBS	72
S103	IAI	58
S104	IAI	65
S107	IAI	35
S107	PR1	76
S104	PR1	68
S107	PR2	60
S106	PR2	43



code	min	max	avg
DBS	72	72	72
IAI	35	65	52.6
PR1	68	76	72
PR2	43	60	51.5

Conditions on AGGREGATE FUNCTIONS

Aggregate function
without conditions

id	code	mark
S103	DBS	72
S103	IAI	58
S104	IAI	65
S107	IAI	35
S107	PR1	76
S104	PR1	68
S107	PR2	60
S106	PR2	43

code	avg
DBS	72
IAI	52.6
PR1	72
PR2	51.5

```
SELECT code,  
       AVG(mark)  
  FROM grade  
  
 GROUP BY code
```



Conditions on AGGREGATE FUNCTIONS

Where conditions operate **BEFORE** the group by.

id	code	mark
S103	DBS	72
S103	IAI	58
S104	IAI	65
S107	IAI	35
S107	PR1	76
S104	PR1	68
S107	PR2	60
S106	PR2	43

id	code	mark
S104	IAI	65
S107	IAI	35
S107	PR1	76
S104	PR1	68
S107	PR2	60
S106	PR2	43

code	avg
IAI	48.5
PR1	72
PR2	51.5

```
SELECT code,  
       AVG(mark)  
  FROM grade  
 WHERE id != 'S103'  
 GROUP BY code
```

Conditions on AGGREGATE FUNCTIONS

New keyword **HAVING** operates **after** the group by.

id	code	mark
S103	DBS	72
S103	IAI	58
S104	IAI	65
S107	IAI	35
S107	PR1	76
S104	PR1	68
S107	PR2	60
S106	PR2	43

id	code	mark
S104	IAI	65
S107	IAI	35
S107	PR1	76
S104	PR1	68
S107	PR2	60
S106	PR2	43

code	avg
IAI	48.5
PR1	72
PR2	51.5

code	avg
PR1	72
PR2	51.5

```
SELECT code,  
       AVG(mark)  
  FROM grade  
 WHERE id != 'S103'  
 GROUP BY code  
 HAVING AVG(mark) > 50
```

We now know a lot...

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68
S104	IAI	65
S106	PR2	43
S107	PR1	76
S107	PR2	60
S107	IAI	35

- **CREATE TABLE** - creates a table
- **UPDATE** - change row(s) in a table
- **INSERT** - add a row to a table
- **DELETE** - remove row(s) from a table
- **SELECT** - return the specified column(s) from a table
- **WHERE** - filters row(s) by condition
- **JOIN** - create new tables of "composite entities"
- **GROUP BY** - creates new entities (rows) by grouping entities and computing specified group (aggregate) facts.
- **HAVING** - filters the result of group by

All operate on a single table.

All return a single table.

Think in tables, entities & attributes...

We now know a lot...

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68
S104	IAI	65
S106	PR2	43
S107	PR1	76
S107	PR2	60
S107	IAI	35

- **CREATE TABLE** - creates a table
- **UPDATE** - change row(s) in a table
- **INSERT** - add a row to a table
- **DELETE** - remove row(s) from a table

- **SELECT** - return the specified column(s) from a table
- **WHERE** - filters row(s) by condition
- **JOIN** - create new tables of "composite entities"
- **GROUP BY** - creates new entities (rows) by grouping entities and computing specified group (aggregate) facts.
- **HAVING** - filters the result of group by

All operate on a single table.

All return a single table.

Think in tables, entities & attributes...

Combining this with joins...

1st we can join two (or more) tables.

We then have **a single table**.

Then we can filter them.

Then we can group them.

Then the required columns are extracted and **aggregates** calculated.

```
SELECT student.id,  
       AVG(grade.mark)  
  FROM student, grade  
 WHERE student.id = grade.id  
 GROUP BY student.id
```

Combining this with joins...

1st we can join two (or more) tables.

We then have **a single table**.

Then we can filter them.

Then we can group them.

Then the required columns are extracted and **aggregates** calculated.

```
SELECT student.id,  
       AVG(grade.mark)  
  FROM student, grade  
 WHERE student.id = grade.id  
 GROUP BY student.id
```

Combining this with joins...

1st we can join two (or more) tables.

We then have **a single table**.

Then we can filter them.

Then we can group them.

Then the required columns are extracted and **aggregates** calculated.

```
SELECT student.id,  
       AVG(grade.mark)  
  FROM student, grade  
 WHERE student.id = grade.id  
 GROUP BY student.id
```

Combining this with joins...

1st we can join two (or more) tables.

We then have **a single table**.

Then we can filter them.

Then we can group them.

Then the required columns are extracted and **aggregates** calculated.

```
SELECT student.id,  
       AVG(grade.mark)  
  FROM student, grade  
 WHERE student.id = grade.id  
 GROUP BY student.id
```

```
SELECT student.id,  
       AVG(grade.mark)  
  FROM student, grade  
 WHERE student.id = grade.id  
 GROUP BY student.id
```

grade

id	code	mark
S103	DBS	72
S103	IAI	58
S105	PR1	68

student

id	first	last
S103	John	Smith
S105	Matt	Black

id	first	last	id	code	mark
S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58
S103	John	Smith	S105	PR1	68
S105	Matt	Black	S103	DBS	72
S105	Matt	Black	S103	IAI	58
S105	Matt	Black	S105	PR1	68



```
SELECT student.id,  
       AVG(grade.mark)  
  FROM student, grade  
 WHERE student.id = grade.id  
 GROUP BY student.id
```

grade

id	code	mark
S103	DBS	72
S103	IAI	58
S105	PR1	68

```
SELECT *
```

```
FROM student, grade
```

```
WHERE (student.id = grade.id)
```



student

id	first	last
S103	John	Smith
S105	Matt	Black

id	first	last	id	code	mark
S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58
S103	John	Smith	S105	PR1	68
S105	Matt	Black	S103	DBS	72
S105	Matt	Black	S103	IAI	58
S105	Matt	Black	S105	PR1	68

```
SELECT *
```

```
FROM student, grade
```

```
WHERE (student.id = grade.id)
```



```
SELECT student.id,  
       AVG(grade.mark)  
  FROM student, grade  
 WHERE student.id = grade.id  
 GROUP BY student.id
```

grade

id	code	mark
S103	DBS	72
S103	IAI	58
S105	PR1	68

```
SELECT *
```

```
FROM student, grade
```

```
WHERE (student.id = grade.id)
```



id	first	last	id	code	mark
S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58
S105	Matt	Black	S105	PR1	68

student

id	first	last
S103	John	Smith
S105	Matt	Black

id	first	last	id	code	mark
S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58
S103	John	Smith	S105	PR1	68
S105	Matt	Black	S103	DBS	72
S105	Matt	Black	S103	IAI	58
S105	Matt	Black	S105	PR1	68

```
SELECT *
```

```
FROM student, grade
```

```
WHERE (student.id = grade.id)
```



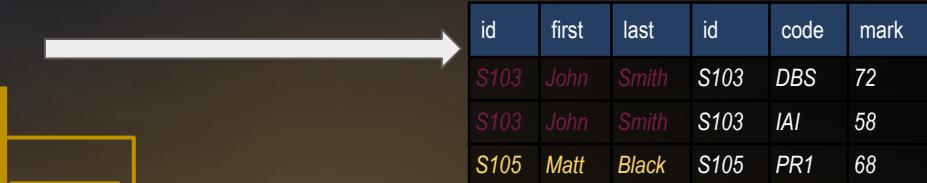
```

SELECT student.id,
       AVG(grade.mark)
  FROM student, grade
 WHERE student.id = grade.id
 GROUP BY student.id

```

grade

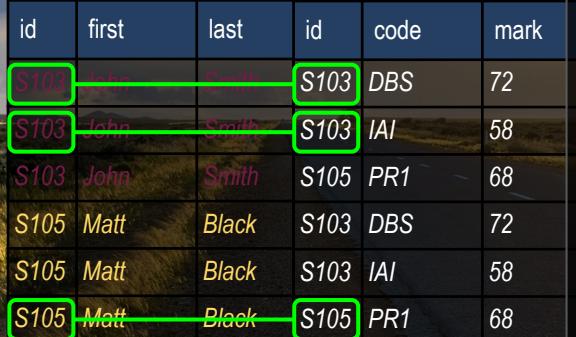
id	code	mark
S103	DBS	72
S103	IAI	58
S105	PR1	68



id	first	last	id	code	mark
S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58
S105	Matt	Black	S105	PR1	68

student

id	first	last
S103	John	Smith
S105	Matt	Black



id	first	last	id	code	mark
S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58
S103	John	Smith	S105	PR1	68
S105	Matt	Black	S103	DBS	72
S105	Matt	Black	S103	IAI	58
S105	Matt	Black	S105	PR1	68



S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58



S105	Matt	Black	S105	PR1	68
------	------	-------	------	-----	----

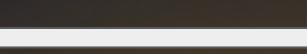
```

SELECT student.id,
       AVG(grade.mark)
FROM student, grade
WHERE student.id = grade.id
GROUP BY student.id

```

grade

id	code	mark
S103	DBS	72
S103	IAI	58
S105	PR1	68



id	first	last	id	code	mark
S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58
S105	Matt	Black	S105	PR1	68



student

id	first	last
S103	John	Smith
S105	Matt	Black

id	first	last	id	code	mark
S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58
S103	John	Smith	S105	PR1	68
S105	Matt	Black	S103	DBS	72
S105	Matt	Black	S103	IAI	58
S105	Matt	Black	S105	PR1	68

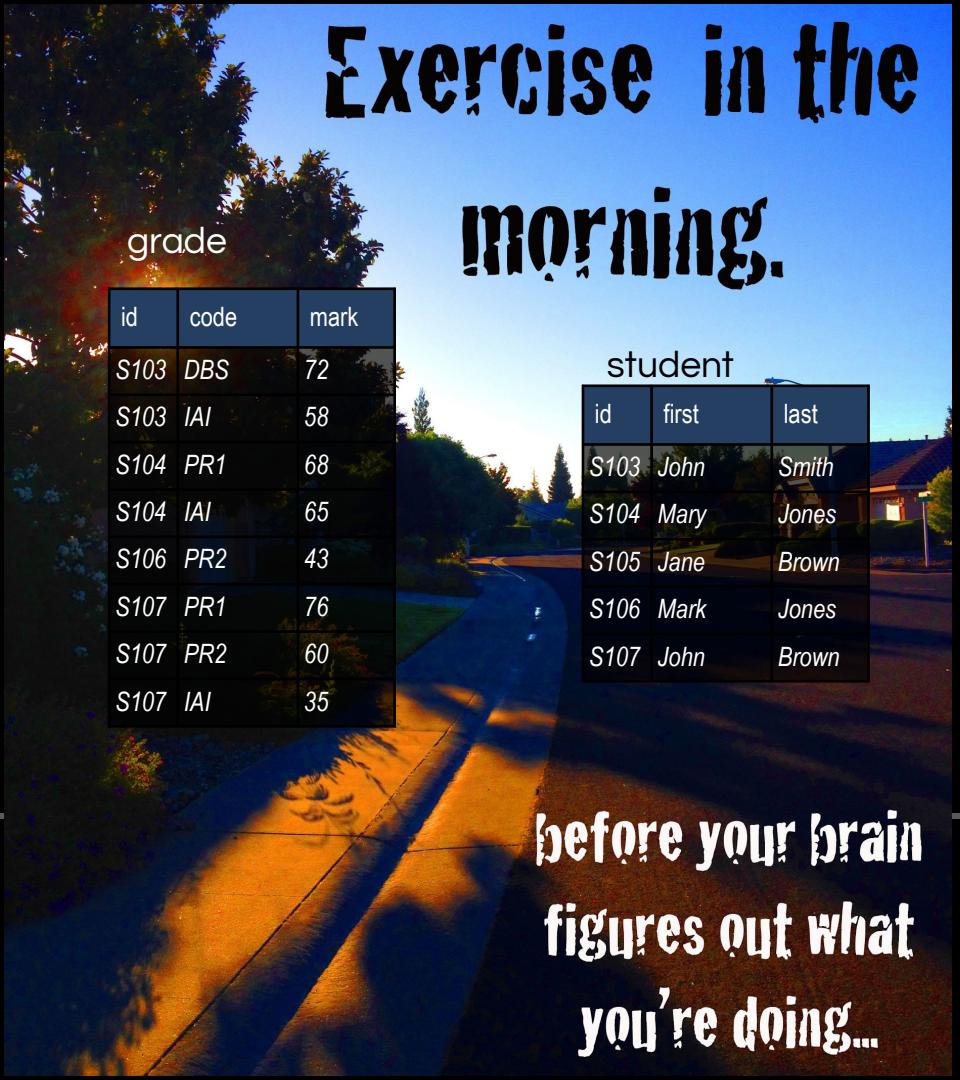


id	first	last	id	code	mark
S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58

id	first	last	id	code	mark
S105	Matt	Black	S105	PR1	68
S105	Matt	Black	S105	PR1	68



id	AVG(mark)
S103	65
S105	68



grade

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68
S104	IAI	65
S106	PR2	43
S107	PR1	76
S107	PR2	60
S107	IAI	35

Exercise in the morning.

student

id	first	last
S103	John	Smith
S104	Mary	Jones
S105	Jane	Brown
S106	Mark	Jones
S107	John	Brown

before your brain
figures out what
you're doing...

```
SELECT id  
FROM student
```

id
S103
S104
S105
S106
S107

```
SELECT *  
FROM grade  
WHERE  
mark >= 60
```

id	code	mark
S103	DBS	72
S104	PR1	68
S104	IAI	65
S107	PR1	76
S107	PR2	60

```
SELECT code, mark  
FROM student, grade  
WHERE first = 'John' and  
last = 'Smith'
```

code	mark
DBS	72
IAI	58

```
SELECT code, AVG(mark)
```

code	mark
DBS	72
IAI	52.667
PR1	72
PR2	53

Recap...

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68
S104	IAI	65
S106	PR2	43
S107	PR1	76
S107	PR2	60
S107	IAI	35

- JOINS via multiple tables in the FROM statement
- Filter to keep only "correct" rows

INNER JOIN

grade

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68

SELECT *

FROM student, grade

WHERE (student.id = grade.id)

student

id	first	last
S103	John	Smith
S105	Jane	Brown

id	first	last	id	code	mark
S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58
S103	John	Smith	S104	PR1	68
S105	Jane	Brown	S103	DBS	72
S105	Jane	Brown	S103	IAI	58
S105	Jane	Brown	S104	PR1	68



INNER JOIN

grade

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68

SELECT *

FROM student, grade

WHERE (student.id = grade.id)



id	first	last	id	code	mark
S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58

student

id	first	last
S103	John	Smith
S105	Jane	Brown

id	first	last	id	code	mark
S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58
S103	John	Smith	S104	PR1	68
S105	Jane	Brown	S103	DBS	72
S105	Jane	Brown	S103	IAI	58
S105	Jane	Brown	S104	PR1	68



INNER JOIN

grade

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68

student

id	first	last
S103	John	Smith
S105	Jane	Brown

id	first	last	id	code	mark
S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58
S103	John	Smith	S104	PR1	68
S105	Jane	Brown	S103	DBS	72
S105	Jane	Brown	S103	IAI	58
S105	Jane	Brown	S104	PR1	68

```
SELECT *  
FROM student, grade  
WHERE (student.id = grade.id)
```

id	first	last	id	code	mark
S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58

What's wrong?

Lost information regarding:

- (1) S105, Jane, Brown
- (2) S104, PR1, 68

Might want to know that join did not match anything!

Select the students that have not recorded any grades.

student

id	first	last
S103	John	Smith
S105	Jane	Brown

```
SELECT *
FROM student, grade
WHERE (student.id = grade.id)
```

grade

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68

Let's consider rows in the original tables.

After the join we have...

**NLAB:***Data at Scale*

student

id	first	last
S103	John	Smith
S105	Jane	Brown

```
SELECT *
FROM student, grade
WHERE (student.id = grade.id)
```

Rows that matched and
were kept (1 or more times)

grade

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68

Let's consider rows in
the original tables.

After the join we
have...



id	first	last	id	code	mark
S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58



NLAB:

Data at Scale

student

id	first	last
S103	John	Smith
S105	Jane	Brown

```
SELECT *
FROM student, grade
WHERE (student.id = grade.id)
```

grade

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68

Let's consider rows in the original tables.

After the join we have...

Rows in **student** that were never matched.



id	first	last	id	code	mark
S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58



NLAB:

Data at Scale

student

id	first	last
S103	John	Smith
S105	Jane	Brown

```
SELECT *
FROM student, grade
WHERE (student.id = grade.id)
```

grade

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68

Let's consider rows in the original tables.

After the join we have...

Rows in
never

grade that were
matched.

id	code	mark
S104	PR1	68

id	first	last
S105	Jane	Brown

id	first	last	id	code	mark
S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58



NLAB:

Data at Scale

student

id	first	last
S103	John	Smith
S105	Jane	Brown

```
SELECT *
FROM student, grade
WHERE (student.id = grade.id)
```

grade

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68

Let's explicitly encode a "did not match anything"

id	first	last
S105	Jane	Brown

id	first	last	id	code	mark
S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58

id	code	mark
S104	PR1	68



NLAB:

Data at Scale

student

id	first	last
S103	John	Smith
S105	Jane	Brown

```
SELECT *
FROM student, grade
WHERE (student.id = grade.id)
```

grade

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68

Let's explicitly encode a "did not match anything"

By matching the line to a set of NULL* values.

* Recall NULL ≈ unknown value

id	first	last	id	code	mark
S105	Jane	Brown	NULL	NULL	NULL

id	first	last	id	code	mark
S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58

id	first	last	id	code	mark
NULL	NULL	NULL	S104	PR1	68



NLAB:

Data at Scale

student

id	first	last
S103	John	Smith
S105	Jane	Brown

```
SELECT *
```

```
FROM student, grade  
WHERE (student.id = grade.id)
```

grade

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68

Back to JOINS

INNER JOIN

Only including matches is called an inner join.

Join all rows to all rows then filter on a match condition.



NLAB:

Data at Scale

student

id	first	last
S103	John	Smith
S105	Jane	Brown

```
SELECT *
FROM student
JOIN grade
ON (student.id = grade.id)
```

grade

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68

New, equivalent syntax!

id	first	last	id	code	mark
S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58

Back to JOINS INNER JOIN

Only including matches is called an inner join.

Join all rows to all rows then filter on a match condition.



NLAB:

Data at Scale

student

id	first	last
S103	John	Smith
S105	Jane	Brown

```
SELECT *
FROM student
JOIN grade
ON (student.id = grade.id)
```

grade

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68

New, equivalent syntax!

id	first	last	id	code	mark
S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58

Back to JOINS INNER JOIN

Only including matches is called an inner join.

id	first	last	id	code	mark
S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58

Join all rows to all rows then filter on a match condition.



NLAB:

Data at Scale

student

id	first	last
S103	John	Smith
S105	Jane	Brown

```
SELECT *
FROM student
LEFT OUTER JOIN grade
ON (student.id = grade.id)
```

grade

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68

Back to JOINS LEFT OUTER JOIN

Join all rows to all rows then filter on a match condition.

THEN add any missing rows in the left table.

id	first	last	id	code	mark
S105	Jane	Brown	NULL	NULL	NULL

id	first	last	id	code	mark
S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58

**NLAB:***Data at Scale*

student

id	first	last
S103	John	Smith
S105	Jane	Brown

```
SELECT *
FROM student
LEFT OUTER JOIN grade
ON (student.id = grade.id)
```

grade

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68

Back to JOINS
LEFT OUTER JOIN

id	first	last	id	code	mark
S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58
S105	Jane	Brown	NULL	NULL	NULL

Join all rows to all rows then
filter on a match condition.

THEN add any
missing rows in the
left table.

id	first	last	id	code	mark
S105	Jane	Brown	NULL	NULL	NULL

id	first	last	id	code	mark
S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58



NLAB:

Data at Scale

student

id	first	last
S103	John	Smith
S105	Jane	Brown

```
SELECT *
FROM student
RIGHT OUTER JOIN grade
ON (student.id = grade.id)
```

grade

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68

Back to JOINS
RIGHT OUTER JOIN

id	first	last	id	code	mark
S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58
NULL	NULL	NULL	S104	PR1	68

Join all rows to all rows then
filter on a match condition.

THEN add any
missing rows in the
right table.

id	first	last	id	code	mark
S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58

id	first	last	id	code	mark
NULL	NULL	NULL	S104	PR1	68



NLAB:

Data at Scale

student

id	first	last
S103	John	Smith
S105	Jane	Brown

```
SELECT *
FROM student
FULL OUTER JOIN grade
ON (student.id = grade.id)
```

grade

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68

Back to JOINS
FULL OUTER JOIN

id	first	last	id	code	mark
S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58
S105	Jane	Brown	NULL	NULL	NULL
NULL	NULL	NULL	S104	PR1	68

Join all rows to all rows then
filter on a match condition.

THEN add any
missing rows in both
the left and right
tables

id	first	last	id	code	mark
S105	Jane	Brown	NULL	NULL	NULL

id	first	last	id	code	mark
S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58

id	first	last	id	code	mark
NULL	NULL	NULL	S104	PR1	68

student

id	first	last
S103	John	Smith
S105	Jane	Brown

```
SELECT *
FROM student
FULL OUTER JOIN grade
ON (student.id = grade.id)
```

If the field is the same in the two tables, we can use the syntax USING instead of ON.

This has the advantage that the output will only include one of the fields student.id and grade.id as we have told it is the same!

```
SELECT *
FROM student
FULL OUTER JOIN grade
USING (id)
```

id	first	last	id	code	mark
S105	Jane	Brown	NULL	NULL	NULL

id	first	last	id	code	mark
S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58

grade

id	code	mark
S103	DBS	72
S103	IAI	58
S104	PR1	68

Back to JOINS FULL OUTER JOIN

id	first	last	id	code	mark
S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58
S105	Jane	Brown	NULL	NULL	NULL
NULL	NULL	NULL	S104	PR1	68

Join all rows to all rows then filter on a match condition.

THEN add any missing rows in both the left and right tables





**Exercise in the
morning.**

**before your brain
figures out what
you're doing...**

